




Deep-learning baseline (BrainNetCNN); LOSO AUC 0.68 < 0.72 (linear FC)

Nonlinearity: $\varphi(x) = \max(0.1x, x)$ (leaky-ReLU)

Objective: $\mathcal{L} = -\sum_k w_k y_k \log \hat{\rho}_k$ (class-weighted cross-entropy)

Optimiser: Adam, lr 10^{-3} , weight decay 10^{-3} , batch 32, 40 epochs · leave-one-site-out